

Libraries Required

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
```

Dataset Import

```
churn_set=pd.read_excel('/content/customer_churn_large_dataset.xlsx')
```

Data Preprocessing

Number of rows and column

```
churn_set.shape
(100000, 8)
```

Showing first five rows of the data set

```
churn_set.head()
```

	CustomerID	Age	Gender	Location	Subscription_Length_Months	\
0	1	63	Male	Los Angeles	17	
1	2	62	Female	New York	1	
2	3	24	Female	Los Angeles	5	
3	4	36	Female	Miami	3	
4	5	46	Female	Miami	19	

	Monthly_Bill	Total_Usage_GB	Churn
0	73.36	236	0
1	48.76	172	0
2	85.47	460	0
3	97.94	297	1
4	58.14	266	0

Check wheather there is null value in dataset or not

```
churn_set.isnull().sum()
```

CustomerID	0
Age	0
Gender	0
Location	0
Subscription_Length_Months	0

```
Monthly_Bill      0
Total_Usage_GB    0
Churn              0
dtype: int64
```

Data Analysis and Visualization

Statistical properties of the dataset

```
churn_set.describe()
```


	CustomerID	Age	Subscription_Length_Months	\
count	100000.000000	100000.000000	100000.000000	
mean	50000.500000	44.027020	12.490100	
std	28867.657797	15.280283	6.926461	
min	1.000000	18.000000	1.000000	
25%	25000.750000	31.000000	6.000000	
50%	50000.500000	44.000000	12.000000	
75%	75000.250000	57.000000	19.000000	
max	100000.000000	70.000000	24.000000	

	Monthly_Bill	Total_Usage_GB	Churn
count	100000.000000	100000.000000	100000.000000
mean	65.053197	274.393650	0.497790
std	20.230696	130.463063	0.499998
min	30.000000	50.000000	0.000000
25%	47.540000	161.000000	0.000000
50%	65.010000	274.000000	0.000000
75%	82.640000	387.000000	1.000000
max	100.000000	500.000000	1.000000


Graph plot (To see which properties effect our label)

```
sns.catplot(x='Churn',data=churn_set,kind='count')
```


```
<seaborn.axisgrid.FacetGrid at 0x785cca2b2c50>
```



```
plot=plt.figure(figsize=(5,5))
sns.barplot(x='Churn',y='Monthly_Bill',data=churn_set)
<Axes: xlabel='Churn', ylabel='Monthly_Bill'>
```



```
plot=plt.figure(figsize=(5,5))
sns.barplot(x='Churn',y='Age',data=churn_set)
<Axes: xlabel='Churn', ylabel='Age'>
```




Correlation

```
correlation=churn_set.corr()  
plt.figure(figsize=(10,10))  
sns.heatmap(correlation,cbar=True,square=True,fmt='.1f',annot=True,annot_kws={'size':8},cmap='Greens')
```

<ipython-input-10-580e05bd6ddd>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
correlation=churn_set.corr()
```

<Axes: >



Data Preprocessing

Removing Label for Churn Analysis

```
X = churn_set.drop('Churn',axis=1)
```

```
print(X)
```

```

      CustomerID  Age  Gender  Location
Subscription_Length_Months \
0              1   63   Male  Los Angeles
17
1              2   62  Female   New York
1

```

```

2          3    24  Female  Los Angeles
5
3          4    36  Female      Miami
3
4          5    46  Female      Miami
19
...      ...    ...      ...      ...
.
99995      99996    33    Male      Houston
23
99996      99997    62  Female      New York
19
99997      99998    64    Male      Chicago
17
99998      99999    51  Female      New York
20
99999      100000    27  Female  Los Angeles
19

```

	Monthly_Bill	Total_Usage_GB
0	73.36	236
1	48.76	172
2	85.47	460
3	97.94	297
4	58.14	266
...
99995	55.13	226
99996	61.65	351
99997	96.11	251
99998	49.25	434
99999	76.57	173

```
[100000 rows x 7 columns]
```

```
churn_set["Gender"].value_counts()
```

```

Female    50216
Male      49784
Name: Gender, dtype: int64

```

Label Binarization

```

Y = churn_set['Churn'].apply(lambda y_value: 1 if y_value>=1 else 0)
print(Y)

```

```

0    0
1    0
2    0
3    1

```

```

4          0
      ..
99995      1
99996      0
99997      1
99998      1
99999      1
Name: Churn, Length: 100000, dtype: int64

```

Model Training and Testing

Splitting data into training and testing

```

X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.2,random_state=1)

print(Y.shape,Y_train.shape,Y_test.shape)

(100000,) (80000,) (20000,)

```

Model: Random Forest Classifier

```

from sklearn.ensemble import RandomForestRegressor

model=RandomForestRegressor()
model.fit(X_train,Y_train)
model.score(X_test,Y_test)

-----
-----
ValueError                                Traceback (most recent call last)
<ipython-input-19-86405439abb2> in <cell line: 2>()
      1 model=RandomForestRegressor()
----> 2 model.fit(X_train,Y_train)
      3 model.score(X_test,Y_test)

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py in fit(self, X, y, sample_weight)
    343         if issparse(y):
    344             raise ValueError("sparse multilabel-indicator for
y is not supported.")
--> 345         X, y = self._validate_data(
    346             X, y, multi_output=True, accept_sparse="csc",
dtype=DTYPE
    347         )

/usr/local/lib/python3.10/dist-packages/sklearn/base.py in _validate_data(self, X, y, reset, validate_separately, **check_params)
    582         y = check_array(y, input_name="y",

```



```

**check_y_params)
    583             else:
--> 584                 X, y = check_X_y(X, y, **check_params)
    585                 out = X, y
    586

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py in
check_X_y(X, y, accept_sparse, accept_large_sparse, dtype, order,
copy, force_all_finite, ensure_2d, allow_nd, multi_output,
ensure_min_samples, ensure_min_features, y_numeric, estimator)
    1104         )
    1105
-> 1106     X = check_array(
    1107         X,
    1108         accept_sparse=accept_sparse,

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py in
check_array(array, accept_sparse, accept_large_sparse, dtype, order,
copy, force_all_finite, ensure_2d, allow_nd, ensure_min_samples,
ensure_min_features, estimator, input_name)
    877         array = xp.astype(array, dtype,
copy=False)
    878             else:
--> 879                 array = _asarray_with_order(array,
order=order, dtype=dtype, xp=xp)
    880                 except ComplexWarning as complex_warning:
    881                     raise ValueError(

/usr/local/lib/python3.10/dist-packages/sklearn/utils/_array_api.py in
_asarray_with_order(array, dtype, order, copy, xp)
    183     if xp.__name__ in {"numpy", "numpy.array_api"}:
    184         # Use NumPy API to support order
--> 185         array = numpy.asarray(array, order=order, dtype=dtype)
    186         return xp.asarray(array, copy=copy)
    187     else:

/usr/local/lib/python3.10/dist-packages/pandas/core/generic.py in
__array__(self, dtype)
    2068
    2069     def __array__(self, dtype: npt.DTypeLike | None = None) ->
np.ndarray:
-> 2070         return np.asarray(self._values, dtype=dtype)
    2071
    2072     def __array_wrap__(

ValueError: could not convert string to float: 'Male'

from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
categorical_features=["Gender", "Location"]

```

```
one_hot=OneHotEncoder()
transformer=ColumnTransformer([("one_hot",
                                one_hot,
                                categorical_features)],
                              remainder='passthrough')

transformed_X=transformer.fit_transform(X)
transformed_X

array([[ 0. ,  1. ,  0. , ..., 17. ,  73.36, 236. ],
       [ 1. ,  0. ,  0. , ...,  1. ,  48.76, 172. ],
       [ 1. ,  0. ,  0. , ...,  5. ,  85.47, 460. ],
       ...,
       [ 0. ,  1. ,  1. , ..., 17. ,  96.11, 251. ],
       [ 1. ,  0. ,  0. , ..., 20. ,  49.25, 434. ],
       [ 1. ,  0. ,  0. , ..., 19. ,  76.57, 173. ]])
```

X.head()

	CustomerID	Age	Gender	Location	Subscription_Length_Months	\
0	1	63	Male	Los Angeles	17	
1	2	62	Female	New York	1	
2	3	24	Female	Los Angeles	5	
3	4	36	Female	Miami	3	
4	5	46	Female	Miami	19	

	Monthly_Bill	Total_Usage_GB
0	73.36	236
1	48.76	172
2	85.47	460
3	97.94	297
4	58.14	266

```
pd.DataFrame(transformed_X)
```

	0	1	2	3	4	5	6	7	8	9	10
11 0 236.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	1.0	63.0	17.0	73.36
1 172.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0	2.0	62.0	1.0	48.76
2 460.0	1.0	0.0	0.0	0.0	1.0	0.0	0.0	3.0	24.0	5.0	85.47
3 297.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0	4.0	36.0	3.0	97.94
4 266.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0	5.0	46.0	19.0	58.14
...
99995 226.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	99996.0	33.0	23.0	55.13

```

99996  1.0  0.0  0.0  0.0  0.0  0.0  1.0  99997.0  62.0  19.0  61.65
351.0
99997  0.0  1.0  1.0  0.0  0.0  0.0  0.0  99998.0  64.0  17.0  96.11
251.0
99998  1.0  0.0  0.0  0.0  0.0  0.0  1.0  99999.0  51.0  20.0  49.25
434.0
99999  1.0  0.0  0.0  0.0  1.0  0.0  0.0  100000.0  27.0  19.0  76.57
173.0

```

```
[100000 rows x 12 columns]
```

```

dummies=pd.get_dummies(churn_set[["Gender","Location"]])
dummies

```

	Gender_Female	Gender_Male	Location_Chicago	Location_Houston
0	0	1	0	0
1	1	0	0	0
2	1	0	0	0
3	1	0	0	0
4	1	0	0	0
...
99995	0	1	0	1
99996	1	0	0	0
99997	0	1	1	0
99998	1	0	0	0
99999	1	0	0	0

	Location_Los Angeles	Location_Miami	Location_New York
0	1	0	0
1	0	0	1
2	1	0	0
3	0	1	0
4	0	1	0
...
99995	0	0	0
99996	0	0	1
99997	0	0	0
99998	0	0	1
99999	1	0	0

```
[100000 rows x 7 columns]
```

```
np.random.seed(24)
```

```
X_train,X_test,Y_train_Train,Y_Test = train_test_split(transformed_X,  
                                                         Y,  
                                                         test_size=0.2)
```

```
model.fit(X_train,Y_train)
```

```
RandomForestRegressor()
```

```
model.score(X_test,Y_test)
```

```
-0.03346709386837565
```

```
X.head()
```

	CustomerID	Age	Gender	Location	Subscription_Length_Months	\
0	1	63	Male	Los Angeles	17	
1	2	62	Female	New York	1	
2	3	24	Female	Los Angeles	5	
3	4	36	Female	Miami	3	
4	5	46	Female	Miami	19	

	Monthly_Bill	Total_Usage_GB
0	73.36	236
1	48.76	172
2	85.47	460
3	97.94	297
4	58.14	266

Predictive System

```
input_data=(2,62,1,0,0,0,0,1,0,1,48.76,172)
```

```
input_data_as_numpy_array = np.asarray(input_data)
```

```
input_data_resaped = input_data_as_numpy_array.reshape(1,-1)
```

```
prediction = model.predict(input_data_resaped)
```

```
print(prediction)
```

```
if(prediction[0]==1):
```

```
    print('Churned')
```

```
else:
```

```
    print('Not Churned')
```

```
[0.53]
```

```
Not Churned
```